

```
import pandas as pd
import numpy as np
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import LabelEncoder, StandardScaler
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score, confusion_matrix, classification_report
import matplotlib.pyplot as plt
import seaborn as sns
```

```
# Load the dataset
df = pd.read_csv('synthetic_ml_dataset.csv')
```

## ✓ *Data Cleaning & Preprocessing*


```
# Drop unnecessary columns
df = df.drop(['Customer_ID', 'Joining_Date', 'Last_Activity_Date'], axis=1)

# Encode categorical variables
label_encoders = {}
for column in ['Gender', 'Employment_Status', 'Marital_Status', 'Region', 'Customer_Churn']:
    le = LabelEncoder()
    df[column] = le.fit_transform(df[column])
    label_encoders[column] = le

# Feature scaling
scaler = StandardScaler()
df[['Annual_Income', 'Credit_Score', 'Account_Balance', 'Loan_Amount']] = scaler.fit_transform(
    df[['Annual_Income', 'Credit_Score', 'Account_Balance', 'Loan_Amount']]
)
```

## ✓ *Feature Selection*

```
# Correlation matrix
plt.figure(figsize=(12, 8))
sns.heatmap(df.corr(), annot=True, cmap='coolwarm')
plt.title('Correlation Matrix')
plt.show()
```



## Model Building

```
# Split the data
X = df.drop('Customer_Churn', axis=1)
y = df['Customer_Churn']
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

```
# Train a Random Forest Classifier
rf = RandomForestClassifier(random_state=42)
rf.fit(X_train, y_train)
```



```
# Predictions
y_pred = rf.predict(X_test)
```

## Model Evaluation


```
# Evaluation metrics
accuracy = accuracy_score(y_test, y_pred)
precision = precision_score(y_test, y_pred)
```

```
recall = recall_score(y_test, y_pred)
```

```
# Print metrics
print(f"Accuracy: {accuracy:.2f}")
print(f"Precision: {precision:.2f}")
print(f"Recall: {recall:.2f}")
print(f"F1 Score: {f1:.2f}")
```

```
↗ Accuracy: 0.68
Precision: 0.00
Recall: 0.00
F1 Score: 0.00
```

```
# Confusion Matrix
conf_matrix = confusion_matrix(y_test, y_pred)
sns.heatmap(conf_matrix, annot=True, fmt='d', cmap='Blues')
plt.title('Confusion Matrix')
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.show()
```



```
# Classification Report
print("\nClassification Report:")
print(classification_report(y_test, y_pred))
```

```
↗
```

Classification Report:				
	precision	recall	f1-score	support
0	0.69	0.99	0.81	138
1	0.00	0.00	0.00	62
accuracy			0.68	200
macro avg	0.34	0.49	0.40	200